

the knee: most of the throughput,
little of the latency

left of the knee:
the accelerator idles,
cost per request is high

throughput

latency

latency per request

throughput per accelerator

1 8 32 128 512

batch size (log scale)

right of the knee: throughput is flat,
the tail grows, the queue waits

