

one model, 8-bit;  
a smaller  
model behind  
a router for easy calls

batch lane: offline  
jobs fill the  
night-time trough

SLOs per class: first token  
and tokens per  
second at p99;  
starvation and utilisation  
on one dashboard

state  
handed over

40 tenants;  
2,000 req/s

gateway: auth,  
tenant token  
budgets, admission

**scheduler:**  
**priority by class,**  
**prefix affinity**

prefill pool:  
compute-bound,  
prefix cache

**decode pool:**  
**continuous batching,**  
**memory-bound**

result cache;  
hit for the  
commonest fifth

prefix cache:  
shared system  
prompts, in memory

stream tokens  
to the tenant

